

✓ R Programming Lab 7

Customer Segmentation and Predictive Analytics Using Machine Learning

Name: Prerna Kale
Roll No.: 24102B2001

```
# Install required packages
packages <- c(
  "readxl", "tidyverse", "lubridate", "cluster", "factoextra",
  "caret", "randomForest", "e1071", "plotly", "pROC", "scales"
)

new_packages <- packages[!(packages %in% installed.packages()[, "Package"])]
if(length(new_packages)) install.packages(new_packages, repos = "https://cloud.r-project.org")

library(readxl)
library(tidyverse)
library(lubridate)
library(cluster)
library(factoextra)
library(caret)
library(randomForest)
library(e1071)
library(plotly)
library(pROC)
library(scales)

set.seed(42)
```

margin

Attaching package: ‘e1071’

The following object is masked from ‘package:ggplot2’:

element

Attaching package: ‘plotly’

The following object is masked from ‘package:ggplot2’:

last_plot

The following object is masked from ‘package:stats’:

filter

The following object is masked from ‘package:graphics’:

layout

Type 'citation("pROC")' for a citation.

Attaching package: ‘pROC’

The following objects are masked from ‘package:stats’:

cov, smooth, var

Attaching package: ‘scales’

The following object is masked from ‘package:purrr’:

discard

The following object is masked from ‘package:readr’:

col_factor

```
data_url <- "https://archive.ics.uci.edu/ml/machine-learning-databases/00352/Online%20Retail.xlsx"
file_path <- "Online_Retail.xlsx"

if(!file.exists(file_path)) {
  download.file(data_url, file_path, mode = "wb")
}

retail <- read_excel(file_path)

cat("Rows:", nrow(retail), "\n")
cat("Columns:", ncol(retail), "\n")
head(retail)
```

Rows: 541909
Columns: 8

A tibble: 6 × 8

InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
<chr>	<chr>	<chr>	<dbl>	<dtm>	<dbl>	<dbl>	<chr>
536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850	United Kingdom
536365	71053	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850	United Kingdom
536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2010-12-01 08:26:00	2.75	17850	United Kingdom
536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	2010-12-01 08:26:00	3.39	17850	United Kingdom
536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	2010-12-01 08:26:00	3.39	17850	United Kingdom
536365	22752	SET 7 BABUSHKA NESTING BOXES	2	2010-12-01 08:26:00	7.65	17850	United Kingdom

```
str(retail)
summary(retail)
```

```
tibble [541,909 × 8] (S3: tbl_df/tbl/data.frame)
 $ InvoiceNo  : chr [1:541909] "536365" "536365" "536365" "536365" ...
 $ StockCode : chr [1:541909] "85123A" "71053" "84406B" "84029G" ...
 $ Description: chr [1:541909] "WHITE HANGING HEART T-LIGHT HOLDER" "WHITE METAL LANTERN" "CREAM CUPID HEARTS COAT HANGER" "KNITTED UNION FL
 $ Quantity   : num [1:541909] 6 6 8 6 6 2 6 6 6 32 ...
 $ InvoiceDate: POSIXct[1:541909], format: "2010-12-01 08:26:00" "2010-12-01 08:26:00" ...
 $ UnitPrice  : num [1:541909] 2.55 3.39 2.75 3.39 3.39 7.65 4.25 1.85 1.85 1.69 ...
 $ CustomerID : num [1:541909] 17850 17850 17850 17850 17850 ...
 $ Country    : chr [1:541909] "United Kingdom" "United Kingdom" "United Kingdom" "United Kingdom" ...

      InvoiceNo      StockCode      Description      Quantity
Length   :541909   Length   :541909   Length   :541909   Min.      :-80995.000
N.unique : 25900    N.unique :  4070    N.unique :  4211    1st Qu.:   1.000
N.blank  :    0     N.blank  :    0     N.blank  :    0     Median :   3.000
Min.nchar:    6     Min.nchar:    1     Min.nchar:    1     Mean  :   9.552
Max.nchar:    7     Max.nchar:   12     Max.nchar:   35     3rd Qu.:  10.000
                        NAs       : 1454     Max.    : 80995.000

      InvoiceDate      UnitPrice      CustomerID
Min.      :2010-12-01 08:26:00   Min.      :-11062.060   Min.      :12346
1st Qu.:2011-03-28 11:34:00     1st Qu.:    1.250     1st Qu.:13953
Median :2011-07-19 17:17:00     Median :    2.080     Median :15152
Mean  :2011-07-04 13:34:57     Mean  :    4.611     Mean  :15288
3rd Qu.:2011-10-19 11:27:00     3rd Qu.:    4.130     3rd Qu.:16791
Max.   :2011-12-09 12:50:00     Max.   : 38970.000     Max.   :18287
                        NAs       :135080

      Country
Length   :541909
N.unique :   38
N.blank  :    0
Min.nchar:    3
Max.nchar:   20
```

```
missing_summary <- data.frame(
  Column = names(retail),
  Missing = sapply(retail, function(x) sum(is.na(x))),
  Missing_Percent = round(sapply(retail, function(x) mean(is.na(x)) * 100), 2)
)

missing_summary
```

A data.frame: 8 × 3

	Column	Missing	Missing_Percent
	<chr>	<int>	<dbl>
InvoiceNo	InvoiceNo	0	0.00
StockCode	StockCode	0	0.00
Description	Description	1454	0.27
Quantity	Quantity	0	0.00
InvoiceDate	InvoiceDate	0	0.00
UnitPrice	UnitPrice	0	0.00
CustomerID	CustomerID	135080	24.93
Country	Country	0	0.00

```
retail_clean <- retail %>%
  mutate(
    InvoiceNo = as.character(InvoiceNo),
    CustomerID = as.character(CustomerID),
    InvoiceDate = as.POSIXct(InvoiceDate),
    Quantity = as.numeric(Quantity),
    UnitPrice = as.numeric(UnitPrice)
  ) %>%
  filter(
    !is.na(CustomerID),
    !grepl("^C", InvoiceNo, ignore.case = TRUE),
    !is.na(Description),
    Quantity > 0,
    UnitPrice > 0
  ) %>%
  mutate(
    Revenue = Quantity * UnitPrice
  )

cat("Original rows:", nrow(retail), "\n")
cat("Rows after preprocessing:", nrow(retail_clean), "\n")
cat("Customers retained:", n_distinct(retail_clean$CustomerID), "\n")
```

Original rows: 541909
Rows after preprocessing: 397884
Customers retained: 4338

```
summary(retail_clean[, c("Quantity", "UnitPrice", "Revenue")])
```

Quantity			UnitPrice			Revenue		
Min.	:	1.00	Min.	:	0.001	Min.	:	1.000e-03
1st Qu.:		2.00	1st Qu.:		1.250	1st Qu.:		4.680e+00
Median	:	6.00	Median	:	1.950	Median	:	1.180e+01
Mean	:	12.99	Mean	:	3.116	Mean	:	2.240e+01
3rd Qu.:		12.00	3rd Qu.:		3.750	3rd Qu.:		1.980e+01
Max.	:	80995.00	Max.	:	8142.750	Max.	:	1.685e+05

```
analysis_date <- max(retail_clean$InvoiceDate) + days(1)

customer_features <- retail_clean %>%
  group_by(CustomerID) %>%
  summarise(
    LastPurchase = max(InvoiceDate),
    Recency = as.numeric(difftime(analysis_date, LastPurchase, units = "days")),
    Frequency = n_distinct(InvoiceNo),
    Monetary = sum(Revenue),
    AvgTransactionValue = Monetary / Frequency,
    TotalQuantity = sum(Quantity),
    UniqueProducts = n_distinct(StockCode),
    ActiveMonths = n_distinct(floor_date(InvoiceDate, "month")),
    .groups = "drop"
  ) %>%
  mutate(
    PurchaseFrequency = Frequency / ActiveMonths
  ) %>%
  select(-LastPurchase)

head(customer_features)
```

A tibble: 6 × 9								
CustomerID	Recency	Frequency	Monetary	AvgTransactionValue	TotalQuantity	UniqueProducts	ActiveMonths	PurchaseFrequency
<chr>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<int>	<int>	<dbl>
12346	326.117361	1	77183.60	77183.6000	74215	1	1	1
12347	2.873611	7	4310.00	615.7143	2458	103	7	1
12348	75.984028	4	1797.24	449.3100	2341	22	4	1
12349	19.124306	1	1757.55	1757.5500	631	73	1	1
12350	310.867361	1	334.40	334.4000	197	17	1	1
12352	36.925694	8	2506.04	313.2550	536	59	4	2

```
customer_features %>%
  summarise(
    Customers = n(),
    AvgRecency = mean(Recency),
    AvgFrequency = mean(Frequency),
    AvgMonetary = mean(Monetary),
    AvgTransactionValue = mean(AvgTransactionValue),
    AvgQuantity = mean(TotalQuantity),
    AvgPurchaseFrequency = mean(PurchaseFrequency)
  )
```

A tibble: 1 × 7

Customers	AvgRecency	AvgFrequency	AvgMonetary	AvgTransactionValue	AvgQuantity	AvgPurchaseFrequency
<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
4338	93.04932	4.272015	2054.266	419.1663	1191.289	1.225594

```
cap_outliers <- function(x) {
  q1 <- quantile(x, 0.25, na.rm = TRUE)
  q3 <- quantile(x, 0.75, na.rm = TRUE)
  iqr_value <- q3 - q1
  lower <- q1 - 1.5 * iqr_value
  upper <- q3 + 1.5 * iqr_value
  x <- pmax(x, lower)
  x <- pmin(x, upper)
  x
}

cluster_features <- c(
  "Recency", "Frequency", "Monetary",
  "AvgTransactionValue", "TotalQuantity",
  "PurchaseFrequency", "UniqueProducts"
)

customer_model <- customer_features %>%
  mutate(across(all_of(cluster_features), cap_outliers))

scaled_features <- scale(customer_model[, cluster_features])
rownames(scaled_features) <- customer_model$CustomerID

cat("Features used for clustering:\n")
print(cluster_features)
```


Features used for clustering:
[1] "Recency" "Frequency" "Monetary"
[4] "AvgTransactionValue" "TotalQuantity" "PurchaseFrequency"
[7] "UniqueProducts"

```
wss <- sapply(2:8, function(k) {
  kmeans(scaled_features, centers = k, nstart = 25)$tot.withinss
})

elbow_data <- data.frame(K = 2:8, WSS = wss)

ggplot(elbow_data, aes(K, WSS)) +
  geom_line() +
  geom_point(size = 3) +
  scale_x_continuous(breaks = 2:8) +
  labs(
    title = "Elbow Method for Customer Segmentation",
    x = "Number of Clusters (K)",
    y = "Within-Cluster Sum of Squares"
  ) +
  theme_minimal()

# The assignment workflow uses 4 clusters for a useful business segmentation.
optimal_k <- 4
cat("Selected number of clusters:", optimal_k, "\n")
```



```
set.seed(42)

kmeans_res <- kmeans(
  scaled_features,
  centers = optimal_k,
  nstart = 50
```

```
)

customer_model$Cluster <- factor(kmeans_res$cluster)

cluster_sizes <- customer_model %>%
  count(Cluster, name = "Customers") %>%
  mutate(Percentage = round(Customers / sum(Customers) * 100, 2))

cluster_sizes
```

A tibble: 4 × 3

Cluster	Customers	Percentage
<fct>	<int>	<dbl>
1	947	21.83
2	1845	42.53
3	771	17.77
4	775	17.87

```
distance_matrix <- dist(scaled_features, method = "euclidean")

hc_res <- hclust(distance_matrix, method = "ward.D2")

plot(
  hc_res,
  labels = FALSE,
  hang = -1,
  main = "Hierarchical Clustering Dendrogram",
  xlab = "Customers",
  sub = ""
)


rect.hclust(hc_res, k = optimal_k, border = 1:optimal_k)

hc_clusters <- cutree(hc_res, k = optimal_k)

customer_model$HC_Cluster <- factor(hc_clusters)

cat("Hierarchical clustering completed with", optimal_k, "clusters.\n")
```

Hierarchical clustering completed with 4 clusters.



```
sil_kmeans <- silhouette(kmeans_res$cluster, distance_matrix)
sil_avg <- mean(sil_kmeans[, "sil_width"])


cat("Average K-Means Silhouette Score:", round(sil_avg, 4), "\n")

fviz_silhouette(sil_kmeans) +
  labs(title = "Silhouette Plot - K-Means Clustering")
```

Average K-Means Silhouette Score: 0.2901

cluster	size	ave.sil.width
1	947	0.41
2	1845	0.30
3	771	0.30
4	775	0.12

Silhouette Plot - K-Means Clustering



```
sil_hc <- silhouette(hc_clusters, distance_matrix)
sil_hc_avg <- mean(sil_hc[, "sil_width"])

comparison <- data.frame(
  Method = c("K-Means", "Hierarchical Clustering"),
  Silhouette_Score = c(sil_avg, sil_hc_avg)
)
```

comparison

A data.frame: 2 × 2

Method	Silhouette_Score
<chr>	<dbl>
K-Means	0.2901121
Hierarchical Clustering	0.2216448


```
pca_res <- prcomp(scaled_features, center = FALSE, scale. = FALSE)

pca_variance <- summary(pca_res)$importance[2, 1:3] * 100
cat("Variance explained by PC1:", round(pca_variance[1], 2), "%\n")
cat("Variance explained by PC2:", round(pca_variance[2], 2), "%\n")
cat("Variance explained by PC3:", round(pca_variance[3], 2), "%\n")

pca_data <- data.frame(
  PC1 = pca_res$x[, 1],
  PC2 = pca_res$x[, 2],
  PC3 = pca_res$x[, 3],
  Cluster = customer_model$Cluster
)

ggplot(pca_data, aes(PC1, PC2, color = Cluster)) +
  geom_point(alpha = 0.65, size = 1.5) +
  labs(
    title = "Customer Segments - PCA 2D Projection",
    x = paste0("PC1 (", round(pca_variance[1], 1), "%)"),
    y = paste0("PC2 (", round(pca_variance[2], 1), "%)"),
  ) +
  theme_minimal()
```


Variance explained by PC1: 59.03 %
Variance explained by PC2: 15.88 %
Variance explained by PC3: 10.98 %
Customer Segments - PCA 2D Projection



```
cluster_profile <- customer_model %>%
  group_by(Cluster) %>%
  summarise(
    Customers = n(),
    Recency = mean(Recency),
    Frequency = mean(Frequency),
    Monetary = mean(Monetary),
    AvgTransactionValue = mean(AvgTransactionValue),
    TotalQuantity = mean(TotalQuantity),
    PurchaseFrequency = mean(PurchaseFrequency),
    UniqueProducts = mean(UniqueProducts),
    .groups = "drop"
  )

cluster_profile
```

A tibble: 4 × 9								
Cluster	Customers	Recency	Frequency	Monetary	AvgTransactionValue	TotalQuantity	PurchaseFrequency	UniqueProducts
<fct>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	947	248.18805	1.374868	322.5573	243.8697	172.3749	1.041887	19.56283
2	1845	49.43259	2.479675	552.1125	236.8786	329.8417	1.099697	35.85203
3	771	23.06724	8.911803	3076.1368	426.8808	1816.2899	1.436010	119.73671
4	775	71.08369	3.067097	1740.5204	578.9053	1030.8408	1.075972	71.71935

```
# Relative profile using standardized feature means
profile_scaled <- customer_model %>%
  group_by(Cluster) %>%
  summarise(across(all_of(cluster_features), mean), .groups = "drop")

profile_scaled
```

A tibble: 4 × 8							
Cluster	Recency	Frequency	Monetary	AvgTransactionValue	TotalQuantity	PurchaseFrequency	UniqueProducts
<fct>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	248.18805	1.374868	322.5573	243.8697	172.3749	1.041887	19.56283
2	49.43259	2.479675	552.1125	236.8786	329.8417	1.099697	35.85203
3	23.06724	8.911803	3076.1368	426.8808	1816.2899	1.436010	119.73671
4	71.08369	3.067097	1740.5204	578.9053	1030.8408	1.075972	71.71935

```
profile_rank <- cluster_profile %>%
  mutate(
    Monetary_Rank = min_rank(desc(Monetary)),
    Frequency_Rank = min_rank(desc(Frequency)),
    Value_Index = Monetary_Rank + Frequency_Rank
  ) %>%
  arrange(Value_Index, desc(Monetary))

high_value_cluster <- as.character(profile_rank$Cluster[1])

cat("High-value cluster identified as:", high_value_cluster, "\n")

customer_model <- customer_model %>%
  mutate(
    IsHighValue = factor(
      ifelse(as.character(Cluster) == high_value_cluster, "Yes", "No"),
```

```
      levels = c("No", "Yes")
    )
  )

table(customer_model$IsHighValue)
```

High-value cluster identified as: 3

No	Yes
3567	771

```
model_data <- customer_model %>%
  select(
    all_of(cluster_features),
    IsHighValue
  )

set.seed(101)

train_index <- createDataPartition(
  model_data$IsHighValue,
  p = 0.80,
  list = FALSE
)

train_data <- model_data[train_index, ]
test_data <- model_data[-train_index, ]

cat("Training rows:", nrow(train_data), "\n")
cat("Testing rows:", nrow(test_data), "\n")
```

Training rows: 3471
Testing rows: 867

```
set.seed(101)

rf_model <- randomForest(
  IsHighValue ~ .,
  data = train_data,
  ntree = 300,
  importance = TRUE
)

rf_pred <- predict(rf_model, test_data)
rf_prob <- predict(rf_model, test_data, type = "prob")[, "Yes"]

rf_cm <- confusionMatrix(
  rf_pred,
  test_data$IsHighValue,
  positive = "Yes"
)

rf_cm
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	709	7
Yes	4	147

Accuracy : 0.9873
95% CI : (0.9774, 0.9936)
No Information Rate : 0.8224
P-Value [Acc > NIR] : <2e-16

Kappa : 0.9562

McNemar's Test P-Value : 0.5465

Sensitivity : 0.9545
Specificity : 0.9944
Pos Pred Value : 0.9735
Neg Pred Value : 0.9902
Prevalence : 0.1776
Detection Rate : 0.1696
Detection Prevalence : 0.1742
Balanced Accuracy : 0.9745

'Positive' Class : Yes

```
set.seed(101)

svm_model <- svm(
  IsHighValue ~ .,
  data = train_data,
  kernel = "radial",
  probability = TRUE
)

svm_pred <- predict(svm_model, test_data, probability = TRUE)
```



```
svm_prob <- attr(
  predict(svm_model, test_data, probability = TRUE),
  "probabilities"
)[, "Yes"]

svm_cm <- confusionMatrix(
  svm_pred,
  test_data$IsHighValue,
  positive = "Yes"
)

svm_cm
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	712	0
Yes	1	154

Accuracy : 0.9988
95% CI : (0.9936, 1)
No Information Rate : 0.8224
P-Value [Acc > NIR] : <2e-16

Kappa : 0.9961

McNemar's Test P-Value : 1

Sensitivity : 1.0000
Specificity : 0.9986
Pos Pred Value : 0.9935
Neg Pred Value : 1.0000
Prevalence : 0.1776
Detection Rate : 0.1776
Detection Prevalence : 0.1788
Balanced Accuracy : 0.9993

'Positive' Class : Yes

```
rf_roc <- roc(
  response = test_data$IsHighValue,
  predictor = rf_prob,
  levels = c("No", "Yes"),
  direction = "<"
)

svm_roc <- roc(
  response = test_data$IsHighValue,
  predictor = svm_prob,
  levels = c("No", "Yes"),
  direction = "<"
)

rf_auc <- as.numeric(auc(rf_roc))
svm_auc <- as.numeric(auc(svm_roc))

performance <- data.frame(
  Model = c("Random Forest", "SVM"),
  Accuracy = c(
    unname(rf_cm$byClass["Sensitivity"] * 0 + rf_cm$overall["Accuracy"]),
    unname(svm_cm$byClass["Sensitivity"] * 0 + svm_cm$overall["Accuracy"])
  ),
  Precision = c(
    unname(rf_cm$byClass["Pos Pred Value"]),
    unname(svm_cm$byClass["Pos Pred Value"])
  ),
  Recall = c(
    unname(rf_cm$byClass["Sensitivity"]),
    unname(svm_cm$byClass["Sensitivity"])
  ),
  F1_Score = c(
    unname(rf_cm$byClass["F1"]),
    unname(svm_cm$byClass["F1"])
  ),
  ROC_AUC = c(rf_auc, svm_auc)
)

performance[, -1] <- round(performance[, -1], 4)
performance
```

A data.frame: 2 × 6

Model	Accuracy	Precision	Recall	F1_Score	ROC_AUC
<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
Random Forest	0.9873	0.9735	0.9545	0.9639	0.9992
SVM	0.9988	0.9935	1.0000	0.9968	1.0000


```
par(mfrow = c(1, 2))

fourfoldplot(
  rf_cm$table,
  color = 1:4,
  main = "Random Forest Confusion Matrix"
)

fourfoldplot(
  svm_cm$table,
  color = 1:4,
  main = "SVM Confusion Matrix"
)

par(mfrow = c(1, 1))
```


Random Forest Confusion Matrix SVM Confusion Matrix



```
plot(
  rf_roc,
  main = "ROC Curves - Random Forest vs SVM",
  lwd = 2
)

plot(
  svm_roc,
  add = TRUE,
  lwd = 2,
  lty = 2
)

legend(
  "bottomright",
  legend = c(
    paste0("Random Forest (AUC = ", round(rf_auc, 3), ")"),
    paste0("SVM (AUC = ", round(svm_auc, 3), ")")
  ),
  lwd = 2,
  lty = c(1, 2)
)
```




```
importance_table <- as.data.frame(importance(rf_model)) %>%
  rownames_to_column("Feature") %>%
  arrange(desc(MeanDecreaseGini))
```

importance_table

A data.frame: 7 × 5

Feature	No	Yes	MeanDecreaseAccuracy	MeanDecreaseGini
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
Frequency	17.809479	30.18684	32.30057	339.03989
Monetary	10.783757	23.00165	22.42468	219.63288
TotalQuantity	9.282757	23.40552	24.77623	171.43787
PurchaseFrequency	25.146771	33.93845	40.07692	128.01151
UniqueProducts	3.440058	36.71664	33.42693	91.11508
Recency	-2.313381	14.41319	11.54177	34.03568
AvgTransactionValue	16.042571	10.90679	20.02757	33.57301

```
varImpPlot(
  rf_model,
  main = "Random Forest Feature Importance"
)
```



```
plot_ly(
  pca_data,
  x = ~PC1,
  y = ~PC2,
  z = ~PC3,
  color = ~Cluster,
  type = "scatter3d",
  mode = "markers",
  marker = list(size = 4)
) %>%
  layout(
    title = "Interactive 3D PCA Customer Segments",
    scene = list(
      xaxis = list(title = "PC1"),
      yaxis = list(title = "PC2"),
      zaxis = list(title = "PC3")
    )
  )
```

```
# Automatically generate a simple segment interpretation table
recommendations <- cluster_profile %>%
  mutate(
    Segment_Type = case_when(
      Monetary >= quantile(Monetary, 0.75) &
        Frequency >= quantile(Frequency, 0.75) &
        Recency <= quantile(Recency, 0.50) ~ "High-value loyal customers",

      Monetary >= quantile(Monetary, 0.75) &
        Frequency < quantile(Frequency, 0.75) ~ "High-spend reactivation opportunity",

      Frequency >= quantile(Frequency, 0.50) &
        Monetary < quantile(Monetary, 0.75) ~ "Frequent moderate-value customers",

      TRUE ~ "Low-engagement customers"
    ),

    Marketing_Strategy = case_when(
      Segment_Type == "High-value loyal customers" ~
        "VIP rewards, loyalty benefits, premium offers and early access",

      Segment_Type == "High-spend reactivation opportunity" ~
        "Personalized reactivation, reminders and targeted offers",

      Segment_Type == "Frequent moderate-value customers" ~
        "Bundles, cross-selling and volume-based promotions",

      TRUE ~
        "Introductory offers, product discovery and low-cost engagement campaigns"
    )
  ) %>%
  select(Cluster, Segment_Type, Marketing_Strategy)

recommendations
```

A tibble: 4 × 3

Cluster	Segment_Type	Marketing_Strategy
---------	--------------	--------------------